SharedLearnable vs. NoVariance: local- and global-scale trajectories

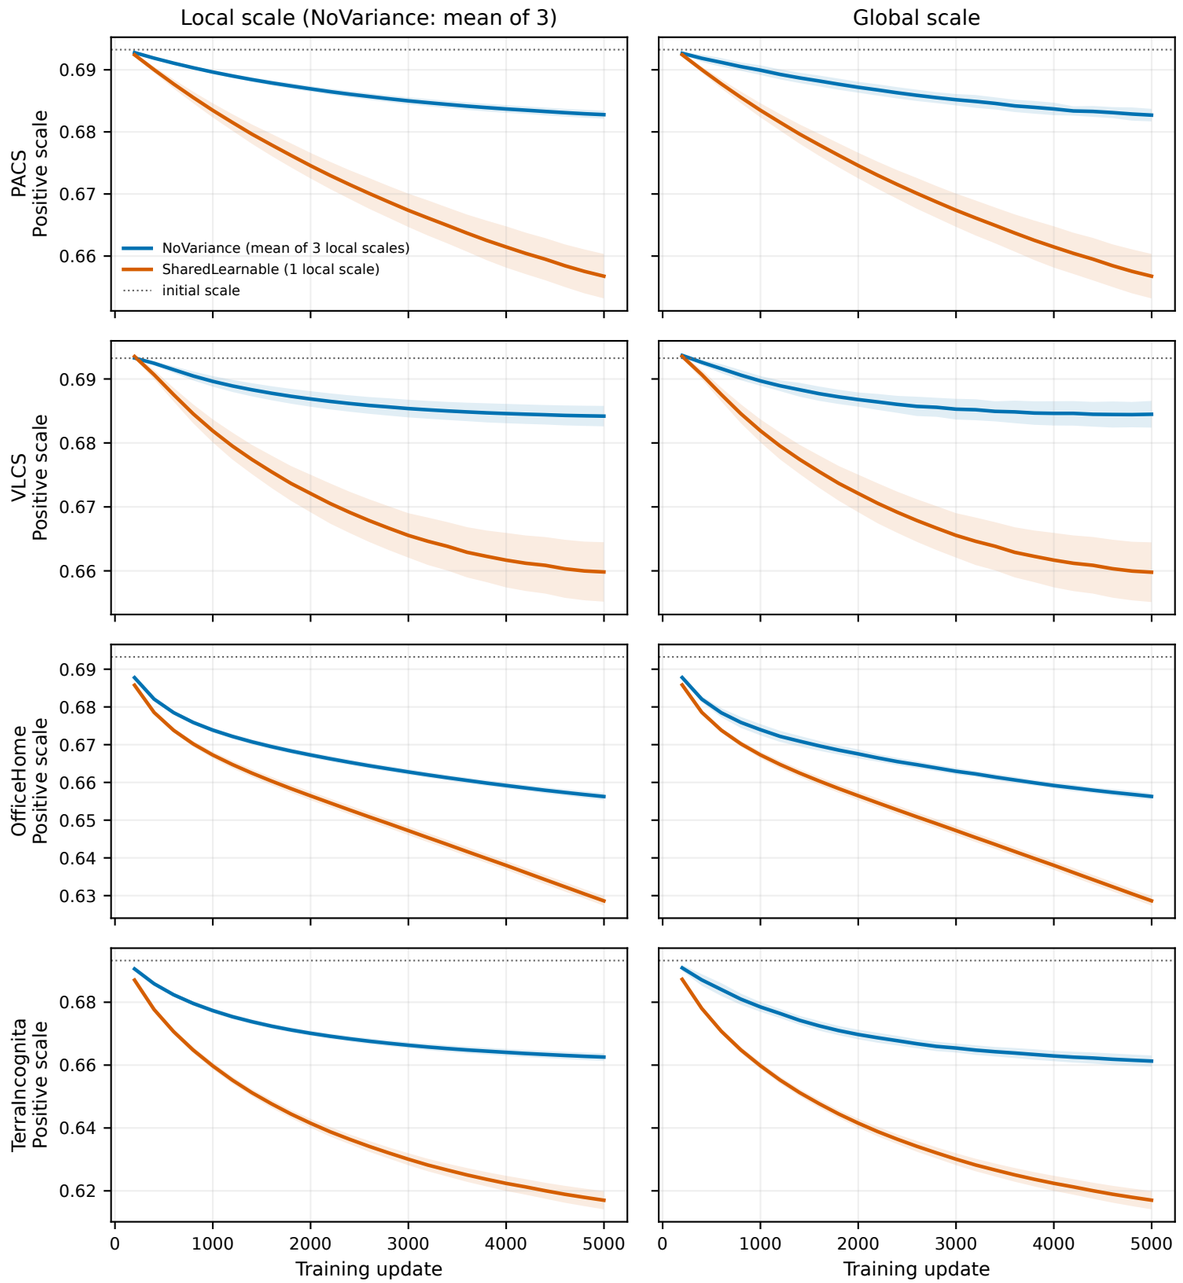

Lines and bands are descriptive mean ± SD over 4 held-out targets × 4 seeds (n=16 runs) per dataset and step. NoVariance's local scale is the per-run mean of its three source-loader-slot values, computed as a summary for this comparison only – individual named-source trajectories appear in Figure 8. SharedLearnable has one local scale by construction.